

Solutions and comments of exercises.

S00

Change the length of the byte slice, from:

```
// TODO: Change a single character in this program so the complete
//         file is read and printed.
```

```
...
b := make([]byte, 11)
...
```

e.g. to 31:

```
...
b := make([]byte, 31)
...
```

Any larger number will do as well. The length of the byte slice is the space that we allow the *Read* method to fill. If this space is too small, we won't be able to read the whole file.

S01

We can shorten these lines

```
// TODO: We don't need to loop manually, there is a helper
//         function for that. Replace the next 10 lines
//         with 5 lines that do the same.
var contents []byte
for {
    b := make([]byte, 8)
    _, err := file.Read(b)
    if err == io.EOF {
        break
    }
    contents = append(contents, b...)
}
fmt.Println(string(contents))
```

by using `io.ReadAll`:

```
b, err := ioutil.ReadAll(file)
if err != nil {
    log.Fatal(err)
}
fmt.Println(string(b))
```

While `io.ReadAll` is useful, it is sometimes overused. Why is that? Often one wants to read something, process it and then write it somewhere else. Imagine a HTTP request body, that we want to read, then preprocess and then maybe write to a file. `io.ReadAll` would consume the *whole data at once*. But for example in the case of a large file upload, there is seldom a reason, why we would need to load the whole file into memory before writing it to disk. There are other ways to solve this problem, which are both more efficient and elegant.

However, `io.ReadAll` is in the standard library and has perfectly fine use cases.

Noteworthy: EOF will not be reported by `ioutil.ReadAll` as the purpose of the method is to consume the reader as a whole:

`ReadAll` reads from `r` until an error or EOF and returns the data it read. A successful call returns `err == nil`, not `err == EOF`. Because `ReadAll` is defined to read from `src` until EOF, it does not treat an EOF from `Read` as an error to be reported.

S02

Use: `io.Copy` and `os.Stdout`.

```
// TODO: Write contents of the file to the standard output,  
//       without using a byte slice (3 lines).  
if _, err := io.Copy(os.Stdout, file); err != nil {  
    log.Fatal(err)  
}
```

The importance of `io.Copy` can hardly be overstated:

`Copy` copies from `src` to `dst` until either EOF is reached on `src` or an error occurs. It returns the number of bytes copied and the first error encountered while copying, if any.

Internally, `io.Copy` uses a buffer in an essential sense:

In computer science, a data buffer (or just buffer) is a region of a physical memory storage used to temporarily store data while it is being *moved from one place to another*.

Everywhere, where readers and writers need to connect, `io.Copy` can be used. As a first example, here we read from a file and write to one of the standard streams.

We will see `io.Copy` over and over again.

S03

Use `os.Stdout` and `os.Stdin`.

```
// TODO: Read input from standard input and pass it to standard output,
//         without using a byte slice (3 lines).
if _, err := io.Copy(os.Stdout, os.Stdin); err != nil {
    log.Fatal(err)
}
```

Here, we have the essence of a filter, namely a program, that works with streams, but does not change the stream at all. One can be reminded of the identify function.

S04

Use `gzip.Reader`.

A `gzip.Reader` is an `io.Reader` that can be read to retrieve uncompressed data from a gzip-format compressed file.

```
// TODO: Read gzip compressed data from standard input and write
//         to standard output, without using a byte slice (7 lines).
r, err := gzip.NewReader(os.Stdin)
if err != nil {
    log.Fatal(err)
}
if _, err := io.Copy(os.Stdout, r); err != nil {
    log.Fatal(err)
}
```

A filter, that decompresses data read from standard input. As soon we get to `io.Copy`, a decompressed stream has the same *shape* as any other type that implements `io.Reader`.

S05

Go comes with an image package in the standard library, which implements a basic 2-D image support.

The fundamental interface is called `Image`.

There is a `Decode` method, that takes a reader and turn it into an `Image`.

In turn, the concrete image subpackages implement an `Encode` method, which takes an `io.Writer` and an `Image` (plus options) as arguments.

```
// TODO: Decode the image and encode it to JPEG, write it
//         to the given writer (5 lines).
// Hint: Utilize methods taking io.Reader or io.Writer
//         in https://golang.org/pkg/image/.
img, _, err := image.Decode(r)
```

```

    if err != nil {
        return err
    }
    return jpeg.Encode(w, img, nil)

```

This snippet takes an arbitrary reader (e.g. standard input) and turns it into an image. The encoding methods are indifferent to the data sink, as long as they implement `io.Writer`.

S06

The package `encoding/json` supports handling streams with `json.Decoder`:

```

// TODO: Unmarshal from standard input into a record struct (4 lines).
var rec record
if err := json.NewDecoder(os.Stdin).Decode(&rec); err != nil {
    log.Fatal(err)
}

```

The decoder takes an `io.Reader` and decodes the read bytes into the given values. This is useful, if you have a possibly large number of values you want to decode, one at a time. The whole stream might not fit into memory at once, but the records, that make up the stream can be processed - one by one.

S07a

Example for a utility reader: A `io.LimitReader` modifies a reader, so that it returns EOF after (at most) a fixed number of bytes.

```

// TODO: Only read the first 27 bytes from standard input (3/6 lines).
if _, err := io.Copy(os.Stdout, io.LimitReader(os.Stdin, 27)); err != nil {
    log.Fatal(err)
}

```

Where can this be useful? Imagine a HTTP response, where the header specifies the content length and you want to limit the reading of the HTTP response body to the number of bytes indicated in the header.

Alternative implementation with a byte slice:

```

// TODO: Only read the first 27 bytes from standard input (3/6 lines).
p := make([]byte, 27)
_, err := os.Stdin.Read(p)
if err != nil {
    log.Fatal(err)
}
fmt.Printf(string(p))

```

Yet another implementation, using `io.CopyN`:

```
// TODO: Only read the first 27 bytes from standard input (3/6 lines).
if _, err := io.CopyN(os.Stdout, os.Stdin, 27); err != nil {
    log.Fatal(err)
}
```

S07b

A `io.SectionReader` wraps seek and read operations. We skip 5 bytes, then read 9 bytes, which should yield the desired string.

We also see that strings can be turned into readers as well.

```
// TODO: Print the string "io.Reader" to stdout (4 lines).
s := io.NewSectionReader(r, 5, 9)
if _, err := io.Copy(os.Stdout, s); err != nil {
    log.Fatal(err)
}
```

Where can this be useful? Imagine a binary file format, that keeps information in various parts of the file and maybe has an index to these sections in a header.

S08

Here, we use `io.ReadFull`, which will read exactly the size of the buffer from the reader.

`ReadFull` reads exactly `len(buf)` bytes from `r` into `buf`. It returns the number of bytes copied and an error if fewer bytes were read. The error is `EOF` only if no bytes were read.

```
// TODO: Read the first 7 bytes of the string into a byte slice,
//         then print to stdout (5 lines).
b := make([]byte, 7)
if _, err := io.ReadFull(r, b); err != nil {
    log.Fatal(err)
}
fmt.Println(string(b))
```

This is a variation of limited reading. Here the limitation is controlled by the size of the byte slice.

S09

We could apply any of the limiting techniques. Here is an example with `io.CopyN`:

```

// TODO: Copy 12 byte from random source into the encoder (3 lines).
if _, err := io.CopyN(encoder, r, 12); err != nil {
    log.Fatal(err)
}

```

If you vary the random seed from call to call, this snippet can serve as a simple version of a password generator.

S10

Another example for `io.Copy`. Here, the destination is a writer, that prettifies tabular data.

```

// TODO: Read tabulated data from standard in
//        and write it to the tabwriter (3 lines).
if _, err := io.Copy(w, os.Stdin); err != nil {
    log.Fatal(err)
}

```

S11

All done.

S12

You can combine any number of readers with `io.MultiReader`.

```

// TODO: Read from these four readers
//        and write to standard output (4 lines).
rs := []io.Reader{
    strings.NewReader("Hello\n"),
    strings.NewReader("Gopher\n"),
    strings.NewReader("World\n"),
    strings.NewReader("! \n"),
}
r := io.MultiReader(rs...)
if _, err := io.Copy(os.Stdout, r); err != nil {
    log.Fatal(err)
}

```

S13

The counterpart to `io.MultiReader` is `io.MultiWriter`. It is similar to the Unix `tee` command.

```

// TODO: Write to both, the file and standard output (4 lines).
w := io.MultiWriter(file, os.Stdout)
if _, err := fmt.Fprintf(w, "SPQR\n"); err != nil {
    log.Fatal(err)
}

```

S14

Fscan belongs to a family of functions, which can be considered the opposite of formatted output: They scan formatted text to yield values.

```

// TODO: Read an int, a float and a string from
//         standard input (3 lines).
if _, err := fmt.Fscan(os.Stdin, &i, &f, &s); err != nil {
    log.Fatal(err)
}

```

S15

Buffers are versatile types. The `bytes.Buffer` is a variable-sized buffer of bytes with `Read` and `Write` methods. You can read a single byte, bytes, runes or a string from it. Writing can be done with similar variety of methods.

```

// TODO: Read one byte at a time from the buffer
//         and print the hex value on stdout (10 lines).
for {
    b, err := buf.ReadByte()
    if err == io.EOF {
        break
    }
    if err != nil {
        log.Fatal(err)
    }
    fmt.Fprintf(os.Stdout, "%x\n", b)
}

```

Here, we read one byte after another. We first check for `io.EOF`, so we can break the loop accordingly. Any other error still needs to be handled. Finally, we use a format verb to format the integer value in base 16, with lower-case letters for a-f.

S16

The `exec.Cmd` struct contains fields for the standard streams, namely `Stdin` of type `io.Reader` and `Stdout` and `Stderr` or type `io.Writer`. Since `bytes.Buffer` is

an `io.Writer` we can connect the standard output of a command directly with a `bytes.Buffer`.

```
// TODO: Stream output of command into the buffer (4 lines).
cmd.Stdout = &buf
if err := cmd.Run(); err != nil {
    log.Fatal(err)
}
```

Imagine, you want to wrap a legacy command line application with a nice Go API. By controlling the input, output and error streams of the application you have basic control over it and you can start parsing and interpreting the command output into Go structures.

S17

A urgent request.

Imagine you get a urgent request to analyze some image data. It's compressed. You need to find the distribution of the “red” values in an image and create a report in form of a pretty table.

This example is short, about 20 lines of code and uses readers and writers all over the place:

- first we read from standard input
- we decompress the data on the fly with a `gzip`
- the image decoding takes a reader
- we use a formatter, that works with a writer
- we use a buffer to temporarily store tab separated values
- we use a `tabwriter` to prettify the data
- we write the final report to standard output

We see, how we can build more complex filters from simple parts.

S18a

We find another important value that is an `io.Reader`: the body of an `http.Response`. It is actually a `ReadCloser`, a type that can be read from and closed.

```
// TODO: Like curl, print the response body to standard output (4 lines).
defer resp.Body.Close()
if _, err := io.Copy(os.Stdout, resp.Body); err != nil {
    log.Fatal(err)
}
```


With the familiar `io.Copy`, we have a simple curl-like program.

While this program would work without the *defer* statement, for serious programs you should always close the response body.

From the `http.Client` documentation:

Caller should close `resp.Body` when done reading from it.

S18b

HTTP is a text based protocol. Instead of using the `http.Get` we craft a similar request ourselves. We write the request string:

```
GET / HTTP/1.0\r\n\r\n
```

to the connection. After that we try to read from the connection and print the result to standard output.

```
// TODO: Send a GET request, read the response and print it to standard output (6 lines)
if _, err := io.WriteString(conn, "GET / HTTP/1.0\r\n\r\n"); err != nil {
    log.Fatal(err)
}
if _, err := io.Copy(os.Stdout, conn); err != nil {
    log.Fatal(err)
}
```

S19

No TODO.

Example for a custom type derived from a file: `atomicfile`. We can write to an atomic file just like any other file:

```
if _, err := io.WriteString(file, "Atomic gopher.\n"); err != nil {
    log.Fatal(err)
}
```

But the semantics are slightly different. Instead of writing to the file directly, all content is written to a temporary file first. The real file is only created (by renaming the temporary file, which is an atomic operation on many operating systems) when the file is closed.

S20

This is our first own `io.Reader` implementation. It does one thing: it says there is nothing to read. Hence the name *Empty*.

```
// TODO: Implement io.Reader interface. Always return EOF (3 lines).
func (r *Empty) Read(p []byte) (n int, err error) {
    return 0, io.EOF
}
```

Even if this is very limited in functionality, you can use this type anywhere, where you can use an `io.Reader`, e.g. in `io.Copy`.

S21

All types implementing `io.Reader` must implement a `Read` method with the exact signature. A typical pattern for implementing custom readers is to keep another reader inside the type, like this:

```
// UpperReader is an uppercase filter.
type UpperReader struct {
    r io.Reader
}
```

Why keep another reader? If we want to uppercase bytes, we have to read the bytes from somewhere. By having the reader available in the type, we make it easy to drop in a filter like this into a processing pipeline. In the example, we can seamlessly connect the `UpperReader` with for example the standard input:

```
...
_, err := io.Copy(os.Stdout, &UpperReader{r: os.Stdin})
...
```

Often, you'll see dedicated constructor for this, like `bufio.NewReader` or `gzip.NewReader`. For our tiny example, we can create the type in a more ad-hoc fashion - but the idea is the same.

```
// TODO: Implement UpperReader, a reader that converts
//      all Unicode letter mapped to their upper case (6 lines).
func (r *UpperReader) Read(p []byte) (n int, err error) {
    n, err = r.r.Read(p)
    if err != nil {
        return
    }
    copy(p, bytes.ToUpper(p))
    return
}
```

The `Read` method implements the core logic. It first reads from the underlying reader. If everything went well, the byte slice `p` will be populated. Now we apply `bytes.ToUpper` to the read bytes and copy them back into the same slice. This works, because

```
bytes.ToUpper(p)
```

is evaluated first, holding the result, which is

a copy of the byte slice [p] with all Unicode letters mapped to their upper case.

This upper-cased version of the byte slice is then copied into the byte slice, that the Read method got as an argument, basically the space we are allowed to populate. Since the mapping from lowercase to uppercase does not change the number of bytes read, we can reuse n as the number of bytes read.

If this looks complicated, please be patient. While this pattern might look unfamiliar at first, it will become familiar the more you are exposed to it and the more you try to implement readers yourself.

S22

A writer that discards everything that is written to it is in a certain way similar to a reader, from which nothing can be read. While it does not seem very useful, Unix users will be reminded of the null device, which actually has reasonable use cases.

```
// TODO: Implement type Discard, which throws away  
//          everything that is written to it (4 lines).  
type Discard struct{}  
  
func (w *Discard) Write(p []byte) (n int, err error) {  
    return len(p), nil  
}
```

We need to implement the Write method with the correct signature. The implementation is simple, since we do not need to do anything at all. We immediately return and report the number of bytes processed and a nil error. Writing to Discard always succeeds.

Why do we return `len(p)` and not just 0?

First, the semantics says that Write should return the number of bytes written:

It returns the number of bytes written from p ($0 \leq n \leq \text{len}(p)$) and any error encountered that caused the write to stop early (io.Writer).

While we do not write anything, we consider all byte to be processed. All bytes will disappear. If we return 0, we actually get an error, namely `io.ErrShortWrite`:

`ErrShortWrite` means that a write accepted fewer bytes than requested but failed to return an explicit error.

S23

Similar to a type, that reads bytes and converts them to upper case, we can implement a writer, that converts all Unicode letter to their upper case. We use another writer, which is where we will write the upper case version to. The implementation is simpler than the corresponding reader, we transform the bytes immediately before we write them to the underlying writer. We pass the return values of the Write method directly back to the caller.

```
// TODO: Implement UpperWriter, a reader that converts  
//      all Unicode letter mapped to their upper case (6 lines).  
type UpperWriter struct {  
    w io.Writer  
}  
  
func (w *UpperWriter) Write(p []byte) (n int, err error) {  
    return w.w.Write(bytes.ToUpper(p))  
}
```

S24a

This counting reader is an example of a reader, that keep track of a metric over time. Here, we just count the number of bytes read in total.

```
// TODO: Implement a reader that counts the total  
//      number of bytes read.  
//      It should have a Count() uint64 method,  
//      that returns the number of bytes read so far.  
//      (12 lines).  
type CountingReader struct {  
    r    io.Reader  
    count uint64  
}  
  
func (r *CountingReader) Read(p []byte) (n int, err error) {  
    n, err = r.r.Read(p)  
    atomic.AddUint64(&r.count, uint64(n))  
    return  
}  
  
// Count returns the total number of bytes read.  
func (r *CountingReader) Count() uint64 {  
    return atomic.LoadUint64(&r.count)  
}
```

The reader uses another reader and keeps a count. In the `Read` method, we increment the count, in this case atomially with the help of the `atomic.AddUint64`.

Finally, we define a public method *Count*, that atomically reads the value with `atomic.LoadUint64` and returns it.

While the number of bytes read is often returned by io-related methods (like `io.Copy`, `io.WriteString`, ...), you can imagine more interesting metrics implemented in this way. For example you could keep track of metrics for the last minute, last five minutes and so on.

S24b

In this example, there is no `TODO`. It is just an example for a more elaborate statistic, that we measure, as data passes through this reader.

The main idea is that - similar to a reader that just count the number of bytes read - we keep the original data untouched and pass it on as is, while inspecting the data and exposing results through additional methods.

Here, we guess the (natural) language of the input. We use a simple trigram-based language model. We could implement this with a simple method as well. The advantage of such a reader based approach would be scalability. It would work even for large files, because we do not need to look at the data all at once. It is ok, if we collect the trigram frequencies chunk by chunk.

S25

In this example, we generate a stream of data. In fact, with a fixed amount of memory we generate endless amount of data. This can be useful in testing scenarios: The time series could be used to test a high frequency trading algorithm or to simulate an internet of things style sensor device emitting a float value every microsecond.

```
$ go run main.go
2017-26-02 20:44:38.901 1.6047
2017-26-02 20:44:38.902 2.2692
2017-26-02 20:44:38.903 1.8446
2017-26-02 20:44:38.904 1.9102
2017-26-02 20:44:38.905 1.8133
2017-26-02 20:44:38.906 1.2980
2017-26-02 20:44:38.907 1.5123
2017-26-02 20:44:38.908 1.1942
2017-26-02 20:44:38.909 0.9112
...
```

The implementation uses an internal buffer to decouple data producer and consumer.

```
// EndlessStream generates a stream of endless data. It uses an internal buffer to
// decouple data production and consumption. When the buffer is empty, we first
// fill the buffer (with linesToFill lines of data), then pass control back to the
// Read method, which then drains the buffer with each call to Read.
type EndlessStream struct {
    buf    bytes.Buffer
    cur    time.Time
    value  float64
}
```

S26

A slow reader. It inserts a short delay between read operations.

- Short asciicast.

S27a

A reader that blacks out text.

```
$ go run main.go
```

```
One morning, when    XX woke from troubled dreams, he found
himself transformed in his bed into a horrible vermin.  He lay on
...
```

We use a `strings.Replacer` to black out words. Since we do not use any internal buffering, we have to take some care not to change the reported number of bytes when replacing words. This is done in the `makeReplacer` helper function. It is also the reason, that we use both and `X` for blacking out the text.

S27b

This is a slightly longer implementation. As usual we use another `io.Reader`. We also add a field of type `time.Duration` to configure the timeout.

```
// TimeoutReader times out, if read takes too long.
// https://github.com/golang/go/wiki/Timeouts
type TimeoutReader struct {
    r        io.Reader
    timeout  time.Duration
}
```

An idiomatic way to implement timeouts can be found in the Go Wiki. The idea is to use a buffered channel, start a goroutine with operation to be performed and then select between two channels: one for the timeout and one for the result of the operation.

We use an auxiliary `readResult` struct for passing the result of a read operation between goroutines.

```
// TODO: Implement a reader that times out after
//       a certain a given timeout (21 lines).
type readResult struct {
    b    []byte
    err error
}

func (r *TimeoutReader) Read(p []byte) (n int, err error) {
    ch := make(chan readResult, 1)

    go func() {
        pp := make([]byte, len(p))
        _, err := r.r.Read(pp)
        ch <- readResult{pp, err}
    }()

    select {
    case <-time.After(r.timeout):
        return 0, ErrTimeout
    case res := <-ch:
        copy(p, res.b)
        return len(p), res.err
    }
}
```

If the read finishes in time, we return the number of bytes read and the error, which might not be nil. If Read takes too long, we return a custom `ErrTimeout`.

The buffered channel allows the goroutine to send a result into: It does not block on send and therefore will be able to finish, even if a timeout occurred.

S28

Example.

The round-robin `MultiReader` is similar to the `io.MultiReader`, but instead of reading each reader until EOF, the round-robin version suspends reading after encountering a newline (or other delimiter) and switches to the next reader. We can read from many readers, and even if the readers are endless, we get a

single stream which consists of a blend of values from all readers (as long, as the delimiter appears regularly in the streams).

The key to the implementation is (again) an internal buffer to decouple production and consumption.

S29

Example.

Combine TimeoutReader and RoundRobinReader. We read in a round robin way from a number of readers. To test the timeout, we will create a *SlowAndFlaky* reader, that is one, that is *sometimes* too slow. This helps us simulate brokenness - a component, that may or may not work.

We set a limit on the number of attempts that should be made to retry another reader. When we start creating readers in main, we create *good* string readers and *bad* SlowAndFlaky readers together. As long as there is enough to read, all is fine. Once all good readers are exhausted and only the bad ones remain, we get closer to the retry limit. If we use enough workers, we will hit the retry limit and the reader will finally fail, saying: *max retries exceeded*.

Short asciicast.

It is often necessary to plan for failure. Not all errors are equal. In this example, we see an example of gradual error handling: we assert the system is ok as long as we can read from *some* reader. We even allow a number of errors. Only when we exceed a certain limit, we finally report an error and fail.

The first fallacy of distributed computing is: the network is reliable. You can imagine many scenarios, where you want gradual error handling: Maybe you don't want to give up on fetching an URL just after a first error. Maybe you have an abstraction over multiple data locations and you do not want to stop, if one is not responding.

The key is to identify errors, that can be recovered from locally and errors that should be passed on.

S30

The bytes package already provides a versatile and performant implementation of a buffer.

However, it is not too complicated to create a simple version of a buffer yourself.

It can be as simple as a byte slice combined with two indices for the current read and write positions:


```

// Buffer is a minimal implementation of a variable size byte slice one can read
// from and write into.
type Buffer struct {
    b    []byte
    i, j int
}

```

When we write to the buffer, we increment `j`, when we read, `i`. A read attempt on an exhausted buffer should return `io.EOF`.

```

// Questions:
//
// (1) This implementation suffers from a small but serious flaw.
//     Can you spot it?
// (2) Can you implement a more efficient version?

```

The flaw is that the current implementation of the buffer does not get rid of memory it does not need any more.

```

// Read reads the current buffer.
func (b *Buffer) Read(p []byte) (n int, err error) {
    sz := len(b.b) - b.i
    if sz > len(p) {
        sz = len(p)
    } else {
        err = io.EOF
    }
    copy(p, b.b[b.i:b.i+sz])
    b.i += sz
    return sz, err
}

```

We read the correct piece of the internal byte slice, but we do not shrink the byte slice after we read from it. We only move pointers. A solution would be to read from the byte slice and to alter the read index *and* to truncate the byte slice.

Snippets S40, S41, S42, S43, S44 are more examples of readers, but they don't contain any exercise.

- S40: Draining a body (duplicates a reader, from the standard library).
- S41: Can we read concurrently from a reader?
- S42: Callbacks (do something of events, such as EOF).
- S43: Like `/dev/zero`.
- S44: A flaky reader that flips bytes.