


Messaging and *job queues* in Go

Talk Structure

1. What is **MOM**, and why?
2. Enterprise MOMs
3. High-speed MOMs
4. **Go** MOMs
5. Wrap-up

Messages Decouple



Message Oriented Middleware


*MOM= any middleware providing
messaging facilities*

Message Oriented Middleware


MOM= any middleware providing messaging facilities

- *asynchronous*: vs “while u wait”
- *decoupled*: mailbox functionality
- *reliable*: store & forward, guarantees
- *scalable*: distributed, load balancing

Request/Reply (RPC)




MOM vs RPC



RPC: [telephone]

- synchronous, tightly coupled processing
- caller is blocked until callee returns
- failure if one endpoint becomes unavailable

MOM vs RPC




RPC: [telephone]

- synchronous, tightly coupled processing
- caller is blocked until callee returns
- failure if one endpoint becomes unavailable

MOM: [mail]

- asynchronous, simple “*fire & forget*” call
- result is queried via message (push or pull)

MOM vs RPC



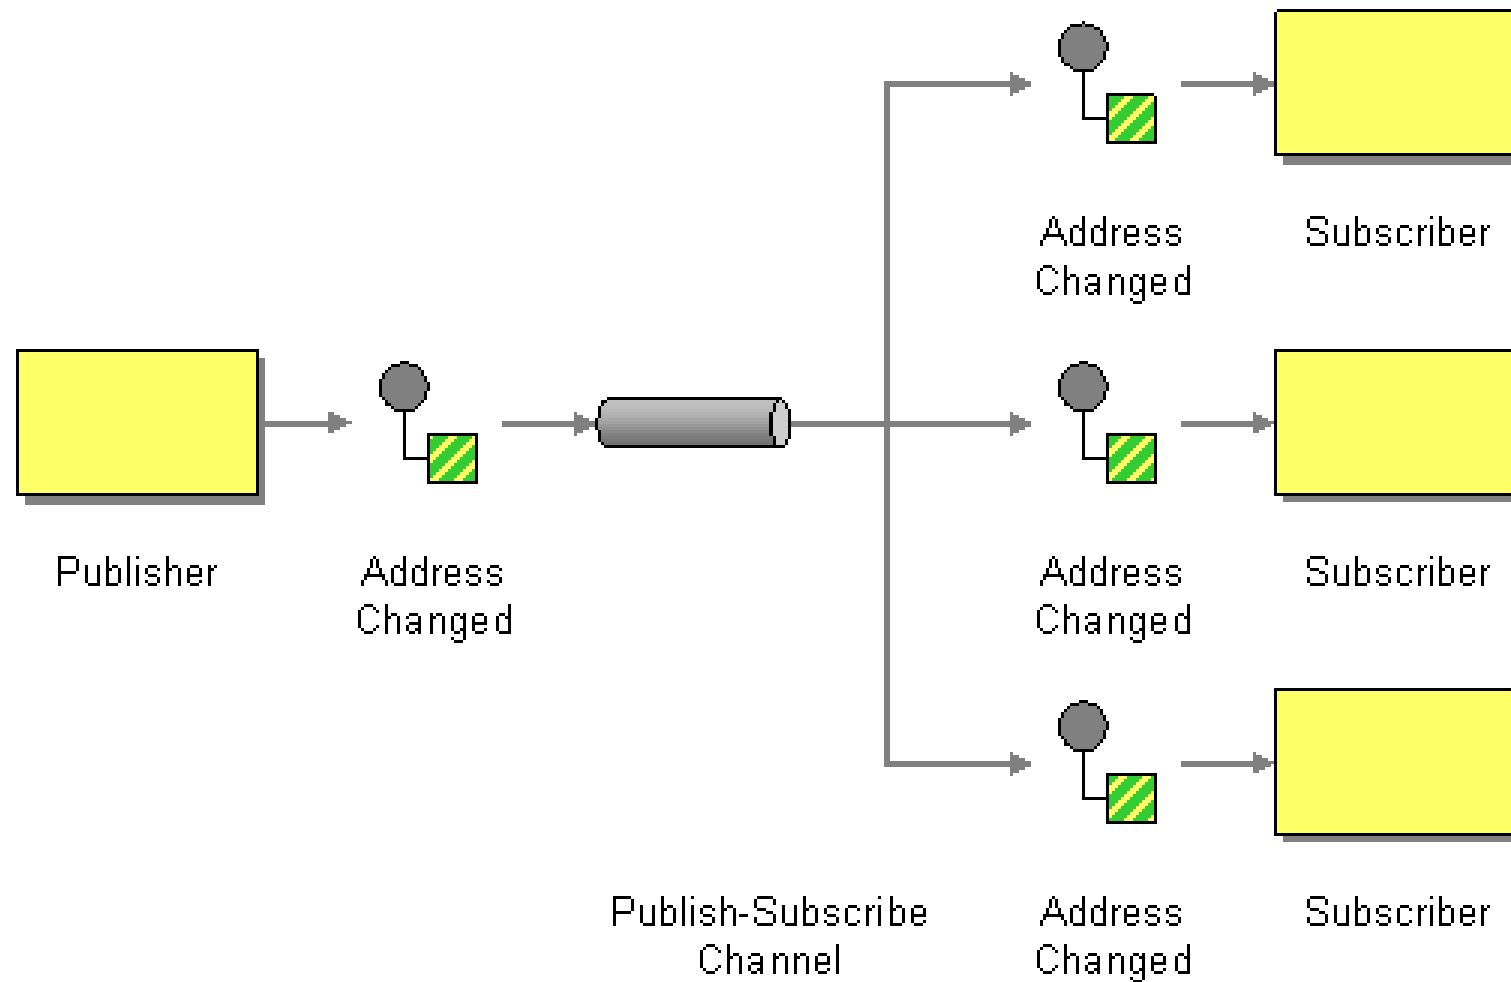
RPC: [telephone]

- synchronous, tightly coupled processing
- caller is blocked until callee returns
- failure if one endpoint becomes unavailable

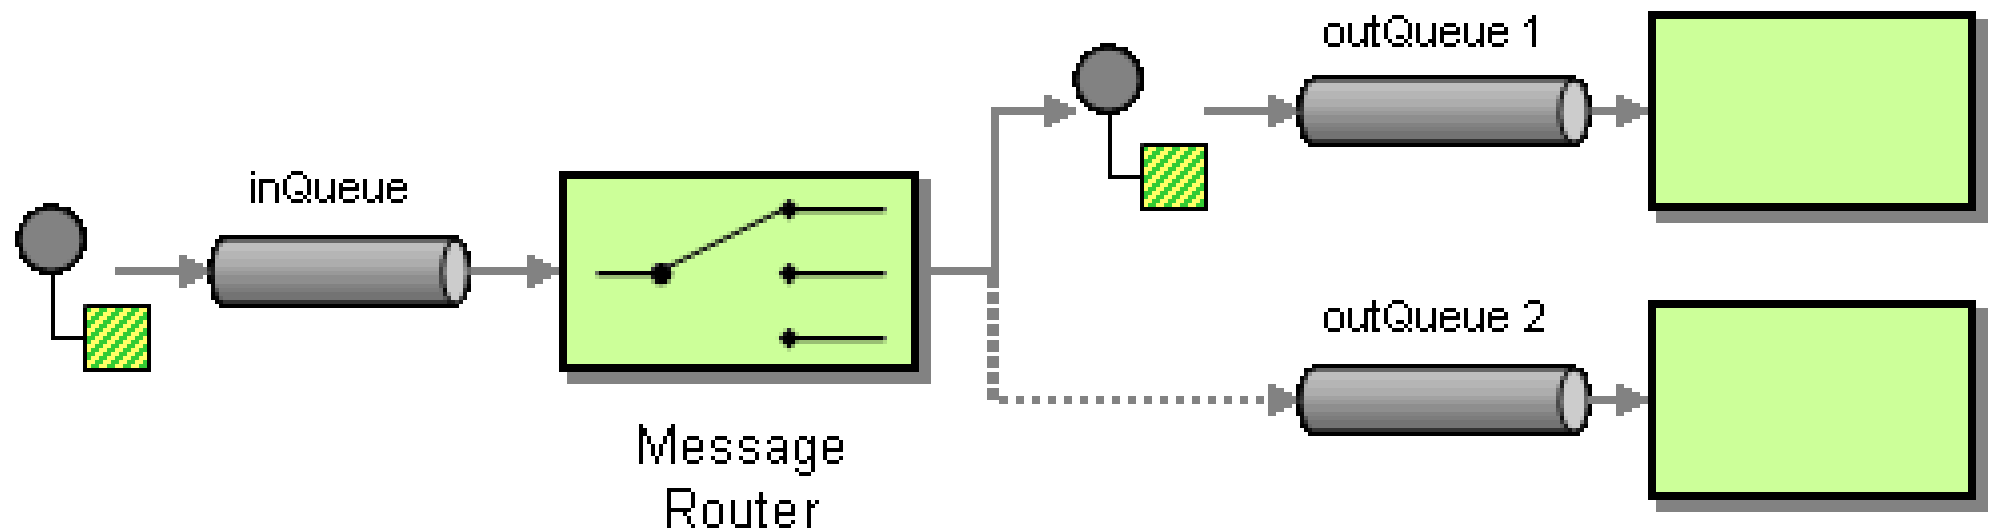
MOM: [mail]

- asynchronous, simple “*fire & forget*” call
- result is queried via message (push or pull)



MOM service: pub/sub queue



MOM service: routing



MOM service: broker



MOM characteristics

- *Queues* decouple:
 - FIFO or priority (job) queues
 - *Message Broker* turns queue into *service*

MOM characteristics

- *Queues* decouple:
 - FIFO or priority (job) queues
 - *Message Broker* turns queue into *service*
- Endpoint models:
 - point-to-point (1:1, peer-to-peer, request/reply)
 - publish/subscribe (m:n)


MOM characteristics

- *Queues* decouple:
 - FIFO or priority (job) queues
 - *Message Broker* turns queue into *service*
- Endpoint models:
 - point-to-point (1:1, peer-to-peer, request/reply)
 - publish/subscribe (m:n)
- QoS guarantees:
 - *at-most-once* (<2, “best effort”)
 - *at-least-once* (>0, needs idempotent receiver)
 - *once-and-once-only* (== 1)

MOM uses and benefits

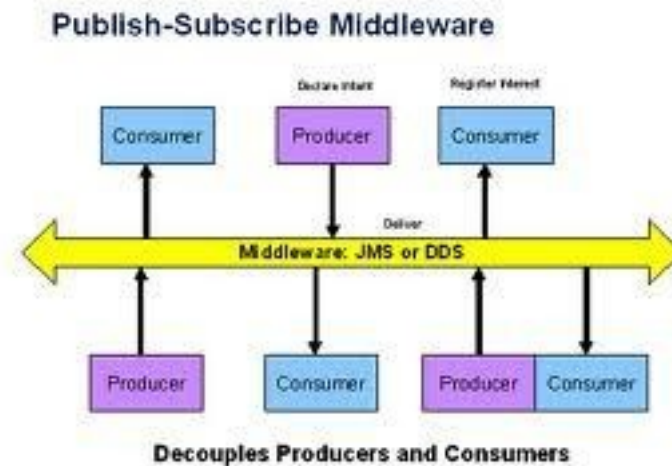
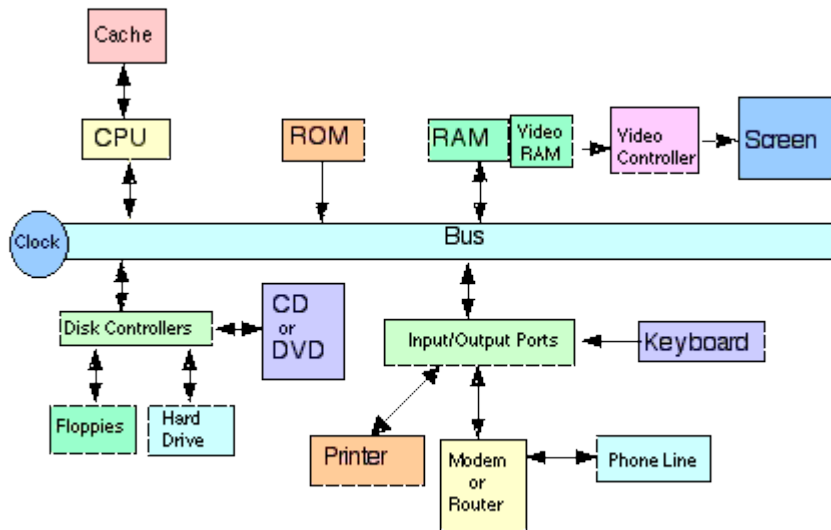
- *Reduce impact of change (EAI)*
 - *tight coupling*: 1 change => *n* other changes
 - *loose coupling*: 1 change in 1 place
- *Simplified application design*
 - messaging component external
 - language-agnostic

2. Enterprise MoMs

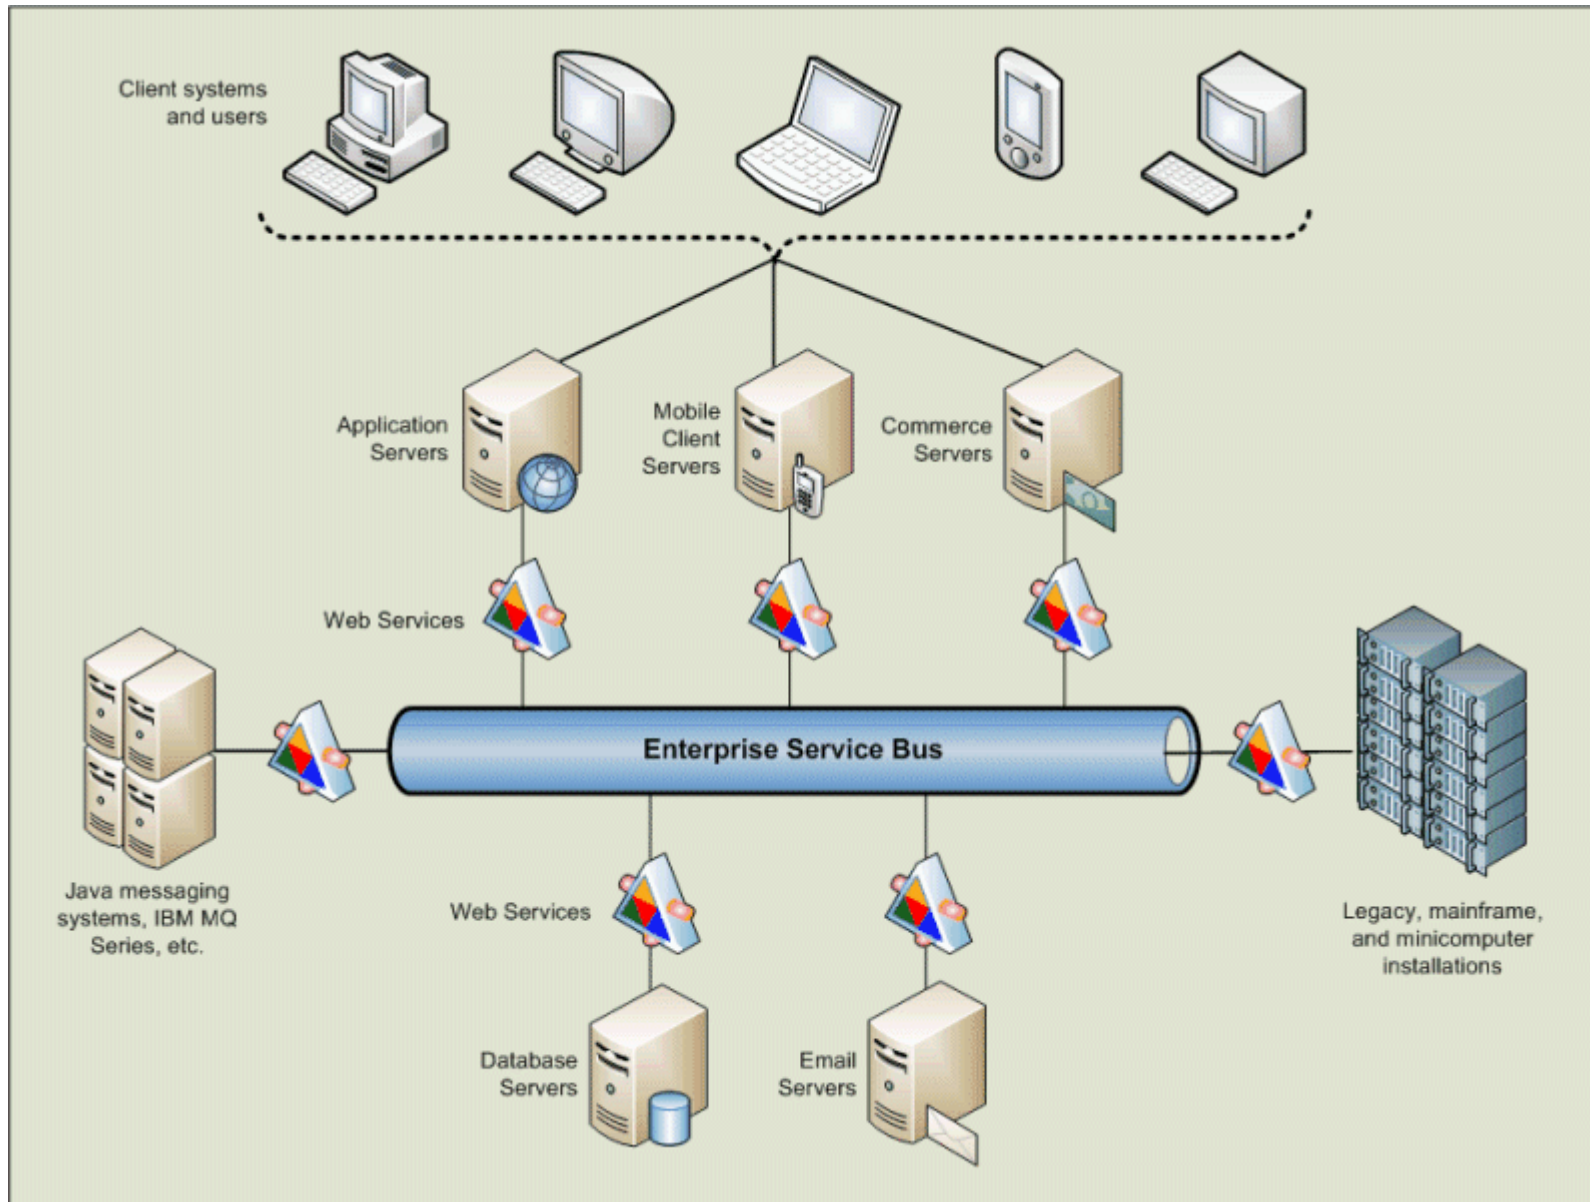


The Information Bus

- *1983: pub/sub origin at TIB*
- Uses *software message bus* for
 - continuous operation (24/7 availability)
 - dynamic evolution (room for extension & change)
 - legacy systems (need to talk to new software)



From TIB to ESB



From TIB to ESB

- 1985: Goldman Sachs
- 1994: Reuters buys TIBCO




From TIB to ESB

- 1985: Goldman Sachs
- 1994: Reuters buys TIBCO



- **1993**: *IBM* MQseries (later Websphere)
- **1997**: *Microsoft* Message Queue (MSMQ)
- **2001**: *Java* Message Service
 - tries to solve vendor-lock in
 - using JDBC-like “*glue API*”
 - JSR-000914



MOM Standardization



AMQP Working Group

Users

JPMorgan
D/Borse
Credit Suisse
Goldman Sachs

TWIST
Novell
29West
Envoy
WSO2



Community

Vendors



Apache
Qpid



iMatix
OpenAMQ



Red Hat
Enterprise
Messaging



Iona
Celtix
AM



RabbitMQ



Cisco
Network

Products

AMQP Defining Features

“enough MOM semantics [...] to meet the needs of most commercial systems”

- *extends* popular JMS semantics

AMQP Defining Features

“enough MOM semantics [...] to meet the needs of most commercial systems”

- *extends* popular JMS semantics
- unlike JMS, defines *wire-level protocol*:
 - *flow-controlled* communication
 - *delivery guarantees* (including exactly-once)
 - *authentication/encryption* based on SASL, TLS
- supports *high-speed trading*

AMQP



- **2003**: JPMorgan Chase (John O'Hara)
- **2004**: iMatix **C** broker (Pieter Hintjens)
- **2005**: Apache Qpid (Java/C++)
- **2006**: RabbitMQ, London (Erlang/OTP)

AMQP



- 2003: JPMorgan Chase (John O'Hara)
- 2004: iMatix C broker (Pieter Hintjens)
- 2005: Apache Qpid (Java/C++)
- 2006: RabbitMQ, London (Erlang/OTP)
- **2008**: AMQP 0.9.1 (“stable”)
- **2011**: AMQP 1.0 (“testing”)
- **2012**: OASIS standard
- **2014**: ISO/IEC 19464

AMQP Basics

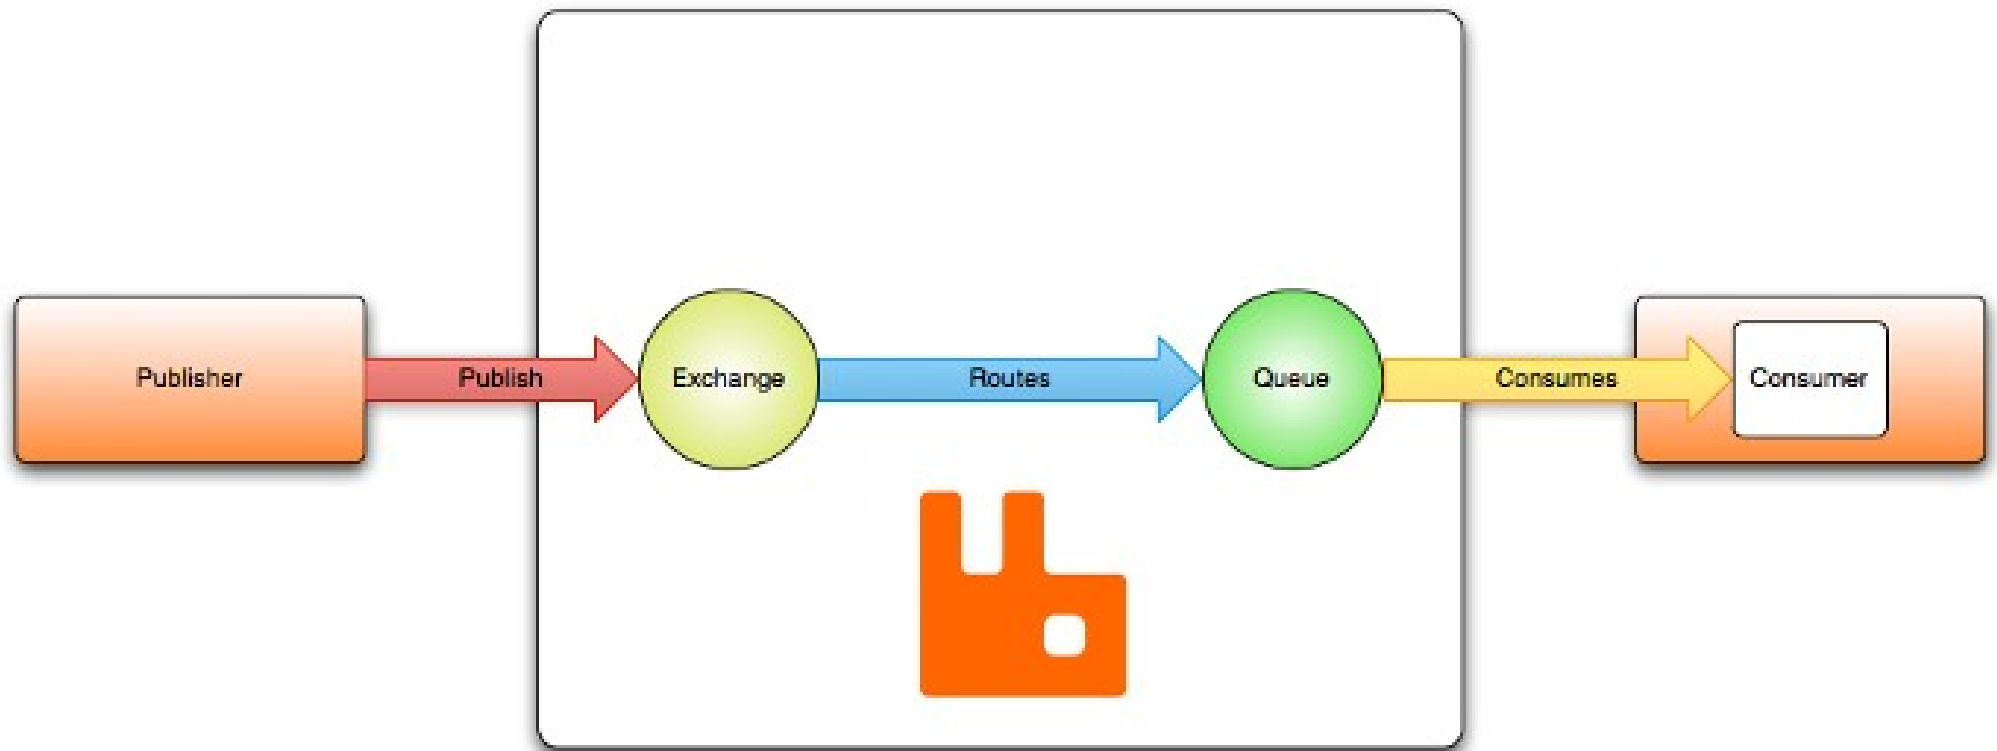
- defines *messaging capabilities*
 - *stores* messages if receiver offline
 - message *tags/topics* used for routing

AMQP Basics

- defines *messaging capabilities*
 - *stores* messages if receiver offline
 - message *tags/topics* used for routing
- **components:**
 - **exchange:** *producer* publishes message to
 - **queue:** stores message for *consumer*
 - **binding:** governs message *routing* from exchange to particular queue(s)

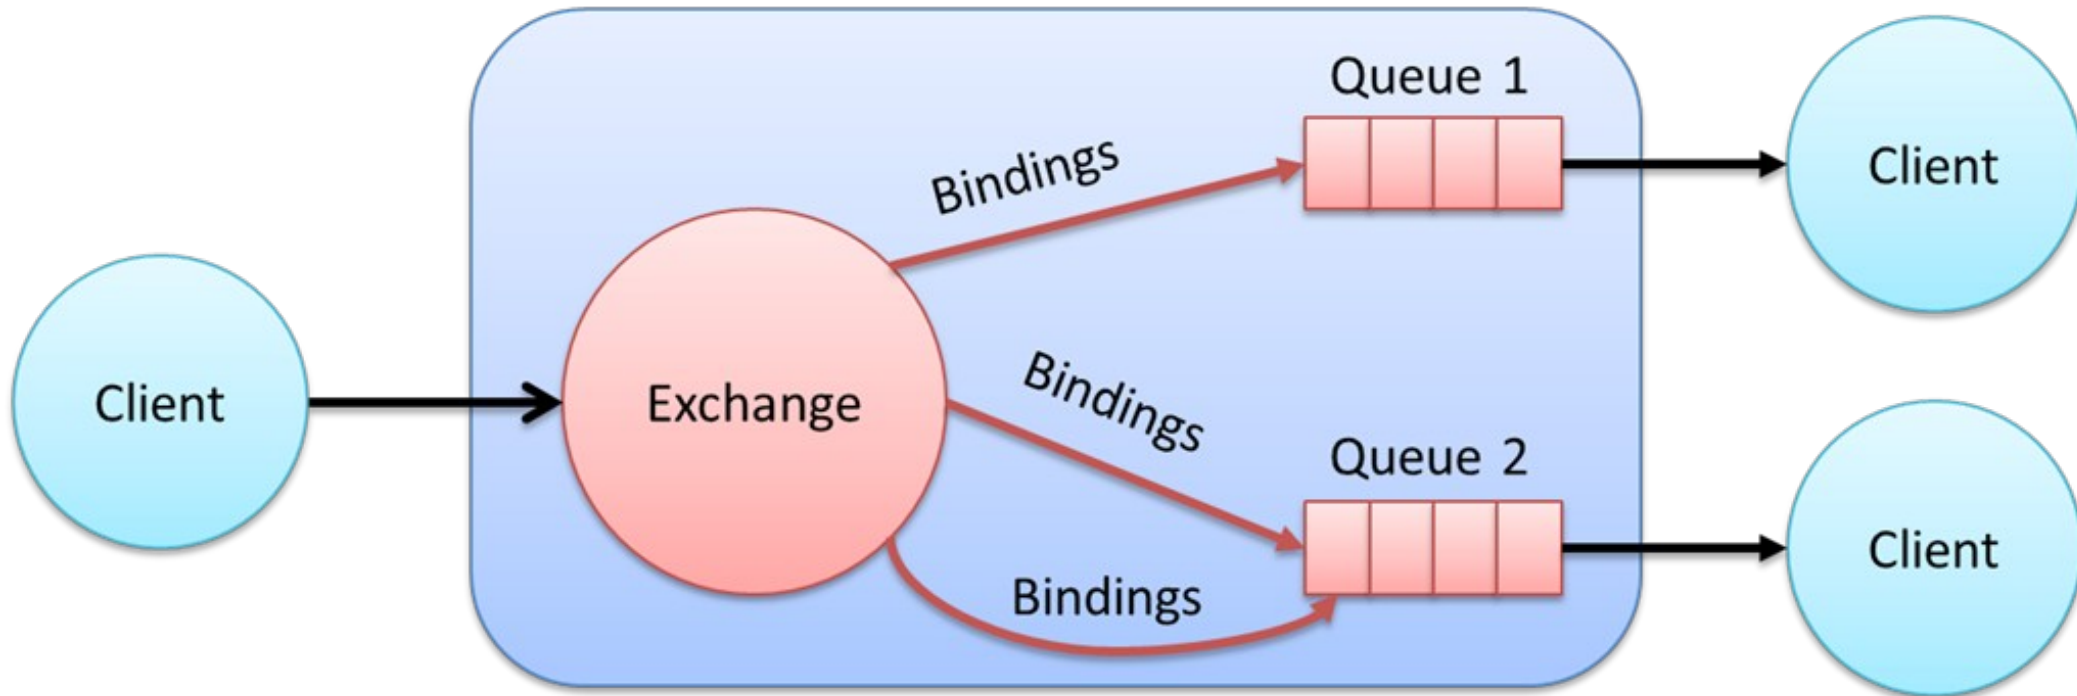
AMQP Basics

"Hello, world" example routing



AMQP Basics

Broker Service





Im in yr serverz,
queueing yr messagez



- implements AMQP 0.9.1
- rock-solid, stable
- based on Erlang/OTP
- *see talk:*

“Evolution of AMQP at Soundcloud”
Sebastian Ohm, Tomás Senart
Erlang User Conference 2013

AMQP to Go

1. Soundcloud RabbitMQ 0.9.1

<https://github.com/streadway/amqp>

2. RabbitMQ Go examples:

<https://github.com/rabbitmq/rabbitmq-tutorials/tree/master/go>

3. Via RabbitMQ HTTP mgmt API:

<https://github.com/michaelklishin/rabbit-hole>

4. Relay AMQP wrapper:

<https://github.com/armon/relay>

5. Apache Qpid Go examples:

<https://github.com/apache/qpid-proton>

Standards after AMQP ...

- AMQP 1.0 spec *drops broker details*
- SOA demand for *reactive systems*:
 - willing to *trade off* speed for consistency and/or reliability
 - 0MQ, LinkedIn, Twitter, Facebook ...

ØMQ

- AMQP is a protocol
- ØMQ is a messaging library:
 - no message queues
 - does not define a broker
 - low-level socket/messaging primitives
 - abstractions for messaging patterns
 - focus on high-performance
- Go clients:
 - <https://github.com/zeromq/goczmq> ZMQ v3
 - <https://github.com/pebbe/zmq4> ZMQ v4
- nanomsg/libmill (go-like concurrency)



Kafka



"By believing passionately in something that still does not exist, we create it. The nonexistent is whatever we have not sufficiently desired."

Franz Kafka

Apache Kafka



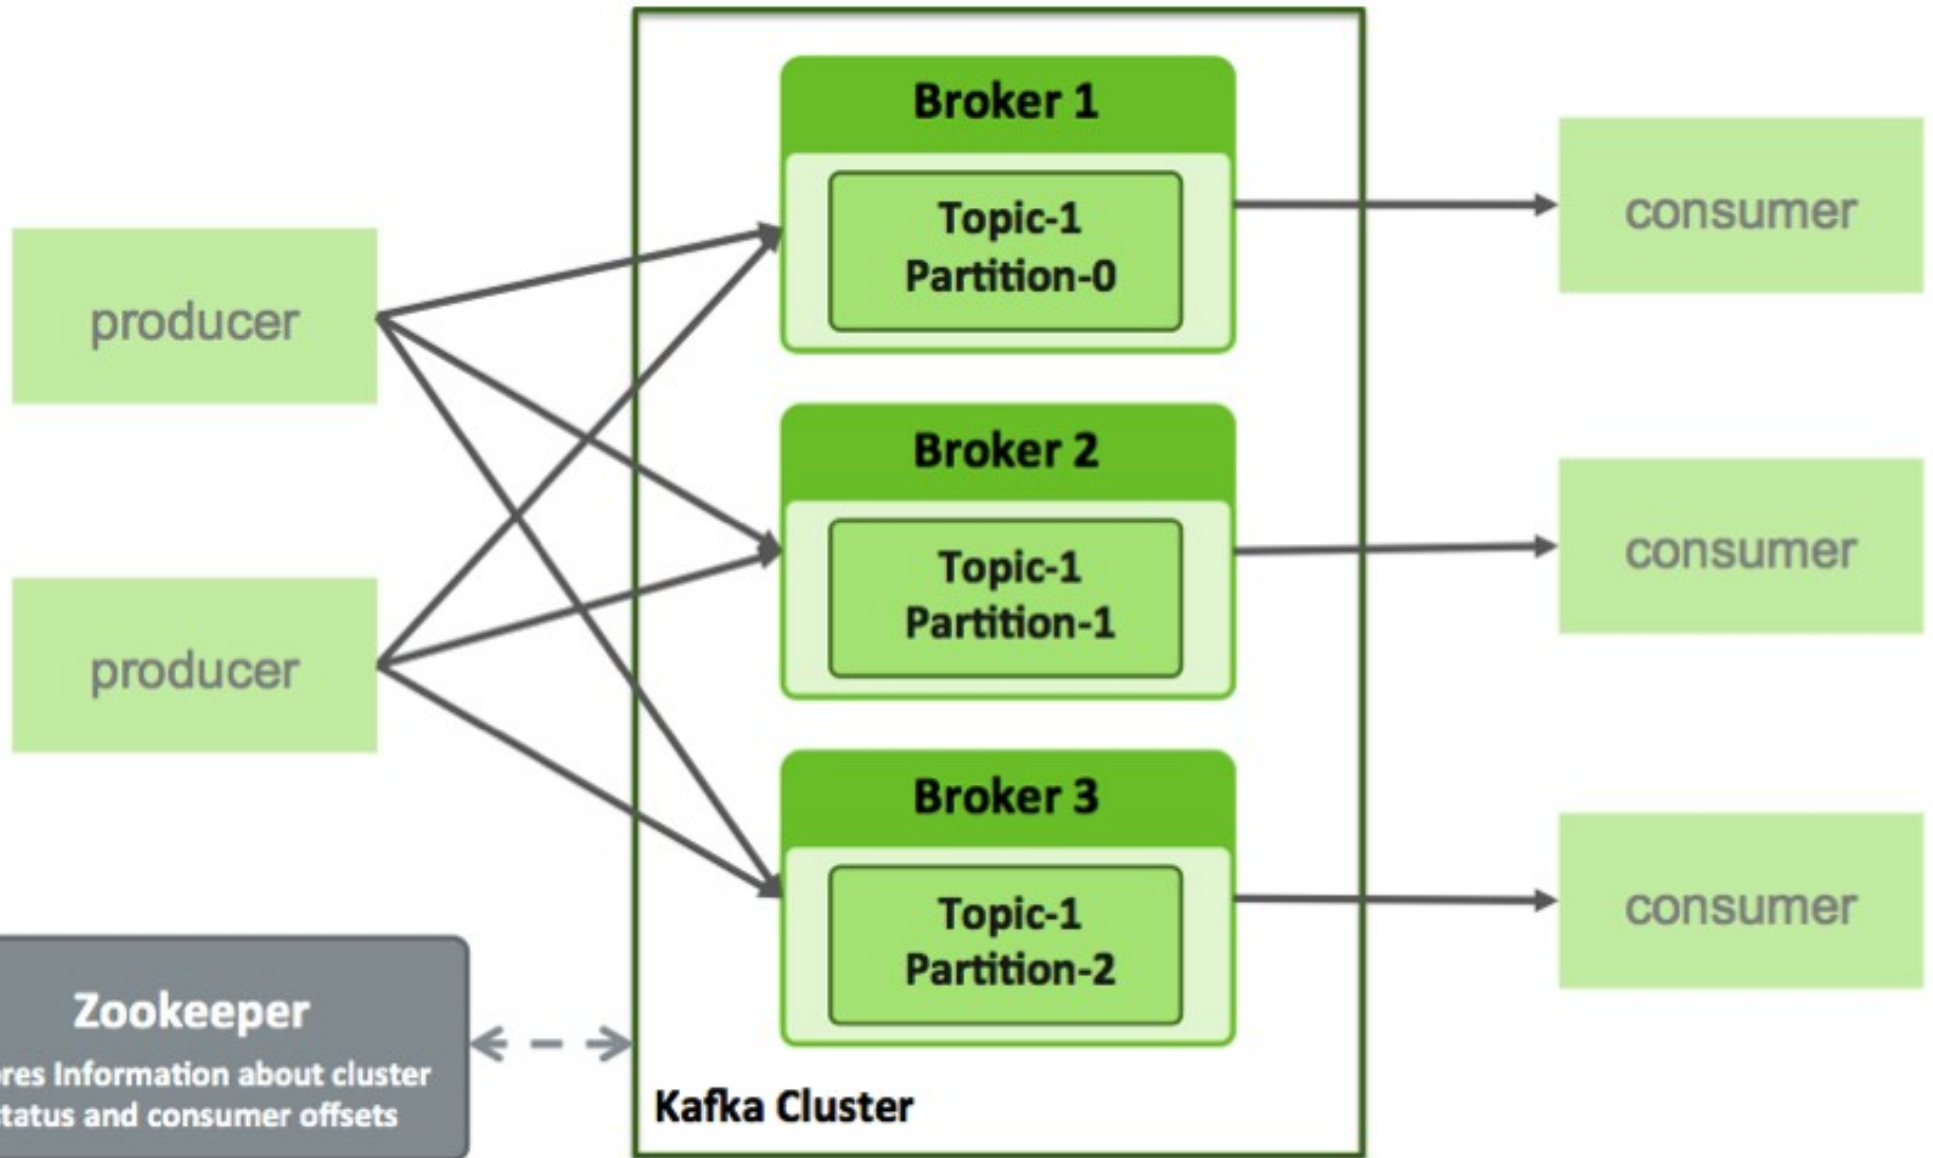
- *broker developed at LinkedIn:*
 - to help with high-speed Hadoop ingest
 - sustain “*firehose throughput*”

Apache Kafka



- *broker developed at LinkedIn:*
 - to help with high-speed Hadoop ingest
 - sustain “*firehose throughput*”
- *unique features:*
 - each *topic* is treated as a *commit log*
 - data partitioning
 - scalable & *low latency*
 - messages are *persisted* on disk and
 - replicated within the cluster

Kafka Architecture



Sarama: Go 4 Kafka



- developed at Shopify

- *Repos:*

<https://github.com/Shopify/sarama>

<https://github.com/eapache>


- *see talk:*

“Complex Concurrency Patterns in Go”

Evan Huus

Golang UK 2015





- *High-performance system* used at Apcera, Baidu, Siemens, vmWare
<https://github.com/nats-io/gnatsd>
- *Trade-offs:*
 - lightweight (server in **Go**, ruby/Go clients)
 - at-most-once delivery
 - no persistence or transactions
- *see talk:*
 - “High Performance Systems in Go”
 - Derek Collison (Apcera)
 - GopherCon 2014



NSQ




- Trade-Offs:
 - for *real-time* distributed messaging
 - horizontal scaling *without brokers*
 - primarily an *in-memory queue*
 - messages not durable
 - *message order not guaranteed*
 - *at-least-once* delivery (idempotent receiver)

<https://github.com/nsqio/nsq>

- *see talk:*
“Spray Some NSQ On It”
Matt Reiferson
GopherCon 2014




Redis & Go



redis

- in-memory store (like memcached)
 - with persistence to disk, more datatypes
 - interesting for realtime, but not reliability


Redis & Go



redis

- in-memory store (like memcached)
 - with persistence to disk, more datatypes
 - interesting for realtime, but not reliability
- *Redis Message Queue in Go:*
<https://github.com/adjust/rmq>

Redis & Go



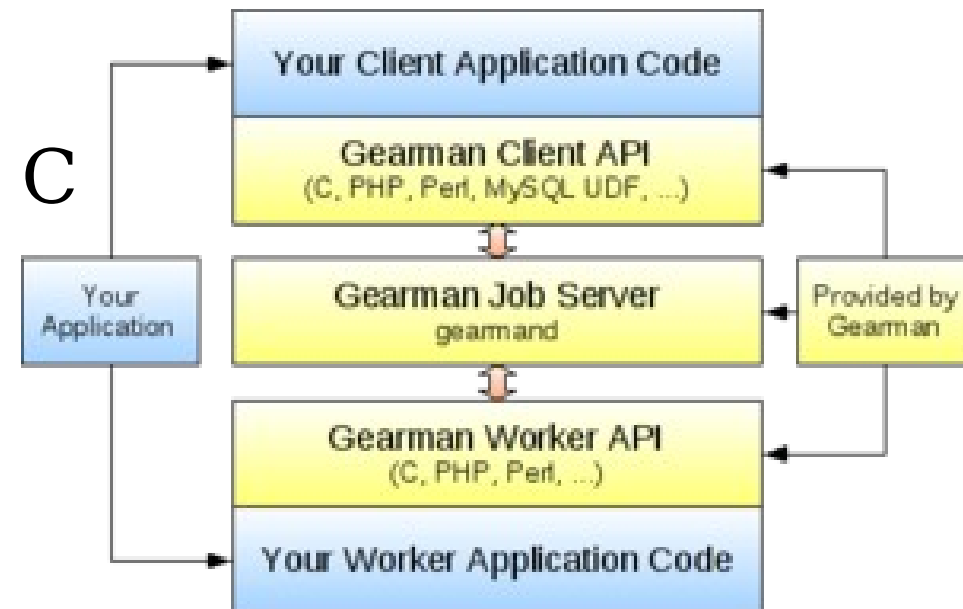
redis

- in-memory store (like memcached)
 - with persistence to disk, more datatypes
 - interesting for realtime, but not reliability
- *Redis Message Queue in Go:*
<https://github.com/adjust/rmq>
- *Github's **resque** for background jobs:*
<http://github.com/resque/resque>
<https://github.com/kavu/go-resque>
- resque compatible **Goworker**:
 - 10..1000 times faster than ruby-based workers
 - <https://www.goworker.org/>



Gearman Job Server

- *Load-balance jobs across workers*
 - conceptually related to map-reduce
 - binary protocol
 - original in Perl, now in C
 - Fitzpatrick: manaGer



- Go Gearman API (client/worker):
<https://github.com/mikespook/gearman-go>

beanstalkd



- fast, simple, in-memory work queue
 - simple, text-based protocol a la memcached
 - 3 *job states* (ready, reserved, buried)
 - transactional
- *Go beanstalk*:
 - <https://github.com/kr/beanstalk>
 - <https://github.com/manveru/gostalk>
 - <https://github.com/iwanbk/gobeanstalk>
 - <https://github.com/99designs/cmdstalk>
 - <https://github.com/nutrun/lentil>

Machinery - Go

- asynchronous *task queue*/job queue
 - to execute distributed jobs in parallel
 - inspired by Python **celery** task queue
- with choice of *broker*:
 - RabbitMQ or
 - Redis

<https://github.com/RichardKnop/machinery>

STOMP

Stomp 

- ***Simple Text Orientated Messaging Protocol***
 - Http-like commands (SEND, SUBSCRIBE, COMMIT, ACK, ABORT, ACK, NACK, ...)
 - <https://stomp.github.io/>
- supported by RabbitMQ, ActiveMQ plugins
- *Go Stomp:*
 - <https://github.com/go-stomp/stomp>
 - <https://github.com/gmallard/stompngo>



In search of Patterns



In search of Patterns

- Go-kit patterns:
 - rate-limiter
 - circuit breaker
 - load-balancer
 - tracing



<https://github.com/go-kit/kit>

- *Talk:*
“Go kit: a toolkit for microservices”
Peter Bourgon
GopherCon 2015, Golang UK 2015

In search of Patterns

- Evan Huus *resiliency patterns*:

- circuit-breaker
- deadline/timeout
- batching
- retrier



<https://github.com/eapache/go-resiliency>

Talk:

“Complex Concurrency Patterns in Go”

Evan Huus

Golang UK, August 2015

Wrap-up

- *breadth of MOM implementations*
- *what about depth?*

Wrap-up

- *breadth of MOM implementations*
- *what about depth?*
- *example: RabbitMQ in Erlang*
 - *complex spec in 5000 lines of code*
 - *solves hard concurrency problems*

Wrap-up

- *breadth of MOM implementations*
- *what about depth?*
- *example: RabbitMQ in Erlang*
 - *complex spec in 5000 lines of code*
 - *solves hard concurrency problems*
- *what are the strengths of Go here?*